

CSE 598: Combinatorial Optimization for Graph Algorithms

Zhengjie Min

September 9, 2026

Abstract

This is a template for TCS (and other kind of) lecture notes.

Contents

I	Flow Algorithms	2
1	Introduction	3
1.1	Menger's Theorem	3
1.2	Nash-Williams Theorem	3
1.3	Edmonds' Theorem	4
1.4	Max-Dicut Packing Min-Dijoin Theorem	4
2	Push Relabel for Max Flow	5
2.1	Setting	5
2.2	Term Definitions	5
2.3	Push-Relabel Algorithm	6
2.3.1	Correctness	6
2.3.2	Time Complexity	7
3	Min Cost Flow	8
3.1	Notations	8
3.2	Optimality Conditions	8
3.3	Algorithm	9
3.3.1	Correctness	10
3.3.2	Time	10
A	Algebra	12
B	Analysis	13

Part I

Flow Algorithms

Chapter 1

Introduction

In this part we'll investigate 4 min-max theorems. In these theorems, a max property of a graph is often proved to be equal to a min property of the graph.

We should be aware that such min-max theorems are not seen everywhere. For example:

Claim.

$$\text{Max Matching} \leq \text{Min Vertex Cover}$$

Proof. Let M be a matching and C be a vertex cover. Then each edge in M must have at least one endpoint in C . Since the edges in M are disjoint, we have that $|M| \leq |C|$. ■

The easiest counterexample to show that the above inequality can be strict is a triangle graph. The maximum matching has size 1, while the minimum vertex cover has size 2.

Here we first list the 4 min-max theorems and prove the $\leq$ direction. Then we will prove the tightness in the next few chapters.

1.1 Menger's Theorem

Claim (Menger's Theorem). In a unweighted graph G , the maximum number of edge-disjoint s - t paths $\leq$ the minimum s - t cut (i.e. the minimum number of edges whose removal disconnects s and t).

Proof. Let $P_1, P_2, \dots, P_k$ be edge-disjoint s - t paths. Let C be an s - t cut. Then each path P_i must contain at least one edge in C . Since the paths are edge-disjoint, we have that $k \leq |C|$. ■

1.2 Nash-Williams Theorem

Claim (Nash-Williams Theorem). The maximum number of edge-disjoint spanning trees $\leq \min_P \lfloor \frac{|E(P)|}{|P|-1} \rfloor$ over all partition P of the graph.

Proof. Let $T_1, T_2, \dots, T_k$ be edge-disjoint spanning trees. Let P be a partition of the graph. Then each spanning tree T_i must contain at least $|P| - 1$ edges in $E(P)$ as a spanning tree must touch all the components of the partition P . Since the spanning trees are edge-disjoint, we have that $k(|P| - 1) \leq |E(P)|$, which implies that $k \leq \lfloor \frac{|E(P)|}{|P|-1} \rfloor$. ■

Note. Intuitively we want to find an equality between the maximum number of disjoint spanning trees and the minimum global cut. However one can find the counterexample of it in C_4 , with the maximum number of disjoint spanning trees being 1, while the minimum global cut is 2.

1.3 Edmonds' Theorem

Definition 1.3.1. A s -**arborescence** in a directed graph is a spanning tree rooted at s such that all vertices are reachable from s .

Claim (Rooted Connectivity Theorem). The maximum number of edge-disjoint s -arborescences $\leq$ the minimum s cut (i.e. the minimum number of edges whose removal disconnects s and any vertex).

We won't be focusing on this theorem. Instead, we want to prove a weaker version in the next few chapters:

Definition 1.3.2. A k **bounded in-degree spanning tree** is a spanning tree packing in a directed graph with:

- k edge-disjoint spanning trees $T_1, T_2, \dots, T_k$.
- Let $T = \bigcup_{i=1}^k T_i$. Then for all $t \neq s$, $\text{indeg}_T(t) = k$.

One can verify that the maximum number of edge-disjoint s -arborescences is at most the maximum number of k bounded in-degree spanning trees.

Claim. The maximum number of k bounded in-degree spanning trees $\leq$ the minimum s cut.

Proof. Let the minimum s cut separate a set X from $V \setminus X$, and $s \in V \setminus X$. As every spanning tree can contain at most $|X| - 1$ edges in the subgraph induced by X , we have

$$|E_T(X)| \leq k(|X| - 1)$$

Also, as every vertex in X have in degree k , we have

$$\sum_{v \in X} \text{indeg}_T(v) = k|X|$$

Combining the two inequalities, we have

$$|E(V \setminus X, X)| = k|X| - |E_T(X)| \geq k|X| - k(|X| - 1) = k$$

which implies that the minimum s cut is at least k . ■

1.4 Max-Dicut Packing Min-Dijoin Theorem

Definition 1.4.1. A **dicut** is the edge set $\delta^{\text{out}}(X)$ for any set X that $\delta^{\text{in}}(X) = \emptyset$.

A **dijoin** is a set of edges that, after contraction, makes the graph strongly connected (i.e. every vertex can reach every other vertex).

Claim (Max-Dicut Packing Min-Dijoin Theorem). The maximum number of edge-disjoint dicuts $\leq$ the minimum dijoin.

Proof. Let the maximum number of edge-disjoint dicuts be k . Let $D_1, D_2, \dots, D_k$ be the edge-disjoint dicuts. Let J be a dijoin. Then each dicut D_i must contain at least one edge in J . Since the dicuts are edge-disjoint, we have that $k \leq |J|$. ■

Chapter 2

Push Relabel for Max Flow

2.1 Setting

Definition 2.1.1 ((s, t) -flow problem). $G = (V, E), U : E \rightarrow \mathbb{R}_{\geq 0}, s, t \in V$. An (s, t) -flow is a function $f : E \rightarrow \mathbb{Z}_{\geq 0}$ such that:

- Capacity: $0 \leq f(e) \leq U(e)$ for all $e \in E$.
- Conservation: $f(\delta^{\text{in}}(v)) = f(\delta^{\text{out}}(v))$ for all $v \neq s, t$.

The value of the flow is $\text{val}(f) = f(\delta^{\text{out}}(s)) - f(\delta^{\text{in}}(s))$.

Definition 2.1.2 ((s, t) -cut). An (s, t) -cut is a partition of V into A and B such that:

- $s \in A$ and $t \in B$.
- The capacity (or size) of the cut is $U(\delta^{\text{out}}(A))$.

Theorem 2.1.1. For $B \geq 0$, there's an algorithm that, in $O(mn^2)$ time, either:

- Finds an (s, t) -flow of value B , or
- Finds an (s, t) -cut of size $< B$.

Corollary 2.1.1. The maximum flow is equal to the minimum cut.

Proof. In last chapter, we have shown that the maximum flow is at most the minimum cut. Now we can use binary search to find the maximum flow. If the maximum flow is B , then we can find an (s, t) -flow of value B and an (s, t) -cut of size $< B + 1$. This implies that the minimum cut is at most B , which implies that the maximum flow is equal to the minimum cut.

Alternatively, if max flow is not equal to the min cut, we can find a value between them and run the algorithm. If the algorithm finds a flow of that value, then the max flow is at least that value, which is a contradiction. If the algorithm finds a cut of that value, then the min cut is at most that value, which is also a contradiction. ■

2.2 Term Definitions

Definition 2.2.1 (Pseudoflow). $f : E \rightarrow \mathbb{Z}_{\geq 0}$ is a **pseudoflow** if it respect the capacity constraints, but not necessarily the conservation constraints.

Definition 2.2.2 (Demand Function). $b : V \rightarrow \mathbb{Z}$ is a **demand function**. f **routes** b if

$$\forall v \in V, f(\delta^{\text{out}}(v)) - f(\delta^{\text{in}}(v)) = b(v).$$

In (s, t) -flow problem, we have $b(s) = B$, $b(t) = -B$ and $b(v) = 0$ for all other vertices.

Definition 2.2.3 (Excess). $\text{ex}_f(v) = b(v) + f(\delta^{\text{in}}(v)) - f(\delta^{\text{out}}(v))$ is the **excess** of v . We call a vertex to be **excess** if $\text{ex}_f(v) > 0$ and **deficit** if $\text{ex}_f(v) < 0$.

Definition 2.2.4 (Residual Graph). The **residual graph** G_f is a directed graph with the same vertex set as G and edge set

$$E_f = \{(u, v) \in E : f(u, v) < U(u, v)\} \cup \{(v, u) : f(u, v) > 0\}.$$

The capacity of an edge (u, v) in G_f is defined as:

$$U_f(u, v) = \begin{cases} U(u, v) - f(u, v) & (u, v) \in E, \\ f(v, u) & (v, u) \in E. \end{cases}$$

Definition 2.2.5 (Height Function). $h : V \rightarrow \mathbb{Z}_{\geq 0}$ is a **height function** if for all $(u, v) \in E_f$, $h(u) \leq h(v) + 1$. Any edge that goes from a higher vertex to a lower vertex is called a **downhill edge**, or **admissible edge**.

Definition 2.2.6 (Active Vertex). A vertex v is **active** if $\text{ex}_f(v) > 0$ and $h(v) < n$.

2.3 Push-Relabel Algorithm

Algorithm 2.1: Push-Relabel Algorithm

```

1 Set  $f(e) = 0$  for all  $e \in E$  and  $h(v) = 0$  for all  $v \in V$ ;
2 while there exists an active vertex  $u$  do
3   if there exists an admissible edge  $(u, v)$  then
4     Push  $\max\{U_f(u, v), \text{ex}_f(u)\}$  units of flow from  $u$  to  $v$ ;
5   else
6      $h(u) \leftarrow h(u) + 1$  (relabel);
```

2.3.1 Correctness

We start with some simple claims. It's generally trivial, so we omit the proofs.

Claim. • f is a pseudoflow at all times.

- Only t will be deficit at all times.
- $h(v) \leq n$.
- h is a valid height function (relabeling only occurs when there are no admissible edges).
- $h(t) = 0$.
- Any excess vertex u has a path to s in the residual graph G_f .

Lemma 2.3.1. Algorithm 2.1 is correct.

Proof. By the end of the algorithm, if there're no excess vertices, then f is a valid flow that routes b . Otherwise, if there exists an excess vertex u , then we have $h(u) = n$.

We denote $V_i = \{v \in V : h(v) = i\}$. As there're $n + 1$ distinct integer heights, $\exists i, V_i = \emptyset$. Fix this i , and let $X = V_{>i} = \{v \in V : h(v) > i\}$.

Then in G_f , there're no edges from X to $V \setminus X$, as for any edge (u, v) with $u \in X$ and $v \in V \setminus X$, we have $h(u) \geq i + 1 \geq h(v) + 2$, which implies that (u, v) is not an admissible downhill edge. Also $s \in X$ as mentioned in the claim above, and $t \notin X$ as $h(t) = 0$.

Then it suffices to prove $U(\delta^{\text{out}}(X)) < B$. We have

$$b(X) - U(\delta^{\text{out}}(X)) = b(X) + f(\delta^{\text{in}}(X)) - f(\delta^{\text{out}}(X)) = \text{ex}_f(X) > 0.$$

As $b(X) \leq B$, we have $U(\delta^{\text{out}}(X)) < B$. ■

2.3.2 Time Complexity

Lemma 2.3.2. Algorithm 2.1 runs in $O(n^2m)$ time.

Proof. First we trivially see that every vertex is relabeled $\leq n$ times.

We denote a saturated push as pushing $U_f(u, v)$ units of flow, and an unsaturated push as pushing $\text{ex}_f(u)$ units of flow.

Fix (u, v) . After a saturated push, it will disappear from the residual graph, and will only reappear after 2 relabels of v (and (v, u) is then an admissible edge). As there're $\leq n$ relabels of v , there're $\leq n/2$ saturated pushes on (u, v) . Thus, there're $\leq nm$ saturated pushes in total.

For unsaturated pushes, we can use a potential function $\Phi = \sum_{v \text{ is active}} h(v)$.

After an unsaturated push from u to v , u will no longer be active, and v may become active. As u is 1 unit higher than v , we have Φ decreases by at least 1.

After a relabel of u , Φ increases by at most 1. There're at most n relabels for each vertex, so Φ increases by at most n^2 due to relabels.

After a saturated push from u to v , v may be active, and Φ increases by at most n . There're at most mn saturated pushes, so Φ increases by at most mn^2 due to saturated pushes.

As Φ is always non-negative, the number of unsaturated pushes is at most $mn^2 + n^2 = O(mn^2)$.

In conclusion, the total number of operations is $O(n^2m)$. ■

Combining the correctness and time complexity, we have proved Theorem 2.1.1.

Chapter 3

Min Cost Flow

Problem 3.0.1 (Min Cost Flow). In a graph with capacities U and costs c , the **cost** of a flow f is defined as

$$\text{cost}(f) = \sum_{e \in E} c(e)f(e)$$

Given integral graph G, b . Either

1. find a flow f that routes b and minimizes $\text{cost}(f)$, or
2. report that no such flow exists.

3.1 Notations

Here we introduce more notations, and update some of the notations from the previous chapter. All other notations remain the same.

Definition 3.1.1 (Net Flow).

$$f^{\text{net}}(v) = f^{\text{in}}(v) - f^{\text{out}}(v)$$

f routes b if $f^{\text{net}}(v) + b(v) = 0$ for all $v \in V$. The excess of f at v is $\text{ex}_f(v) = b(v) - f^{\text{net}}(v)$.

Definition 3.1.2 (Residual Graph). Fix G, f . The **residual graph** $G_f = (V, E_f)$ where $\forall e = (u, v)$:

- If $f(e) < U(e)$, then $(u, v) \in E_f$ with capacity $U_f(u, v) = U(u, v) - f(u, v)$ and cost $c_f(u, v) = c(u, v)$.
- If $f(e) > 0$, then $(v, u) \in E_f$ with capacity $U_f(v, u) = f(u, v)$ and cost $c_f(v, u) = -c(u, v)$ (as we intend to reverse such flow).

Definition 3.1.3 (Circulation). A flow f is a **circulation** if it routes $b = 0$. It might be many cycles added together, in this case:

$$f = \sum_i \alpha_i \mathbb{1}_{C_i}$$
$$\text{cost}(f) = \sum_i \alpha_i \text{cost}(C_i)$$

where C_i is a cycle and α_i is the amount of flow on that cycle. It's trivial that we can compute the cost of a circulation in $O(m^2)$ time.

3.2 Optimality Conditions

Lemma 3.2.1 (OPT certificate). Let f route b . Then f has minimum cost if and only if there is no negative cost cycle in G_f .

Proof. If there is a negative cost cycle C in G_f , then we can augment flow along C to get a new flow f' that routes b and has lower cost than f .

Conversely, if f is not optimal, then there exists a flow f' that routes b and has lower cost than f . Let $g = f' - f$. First we claim g is respecting capacity of G_f , as we can first revert G_f to G by applying $-f$, and then find that f' is respecting capacity of G .

Also, g itself is a circulation (as b cancels out), and $\text{cost}(g) = \text{cost}(f') - \text{cost}(f) < 0$. Then we can decompose g into cycles, and at least one of them must have negative cost. ■

Definition 3.2.1 (Reduced Cost). Fix height $h : V \rightarrow \mathbb{R}$. For a residual edge $a = (u, v) \in E_f$, the **reduced cost** of a is defined as

$$c_{f,h}(a) = c_f(a) - h(u) + h(v)$$

Similar to previous chapter, c_f can be called as downhill length, and $c_{f,h}$ as downhill slack.

In a path P , $c_{f,h}(P) = c_f(P) - h(u) + h(v)$ where u is the start of P and v is the end of P .

In a cycle C , $c_{f,h}(C) = c_f(C)$.

Definition 3.2.2 (ϵ -optimality). (f, h) is ϵ -optimal if $c_{f,h}(a) \geq -\epsilon$ for all $a \in E_f$.

Lemma 3.2.2. Let f be a feasible flow (a.k.a. routes b).

If (f, h) is ϵ -optimal when $\epsilon < 1/n$, then there's no negative cycle in G_f (which implies f is optimal).

Proof. Suppose there's a negative cycle Q in G_f . Then $c_f(Q) \leq -1$.

However as (f, h) is ϵ -optimal, we have $c_f(Q) = c_{f,h}(Q) \geq |Q| \cdot (-\epsilon) \geq -n \cdot \epsilon > -1 \Rightarrow \text{contradiction}$. ■

3.3 Algorithm

Algorithm 3.1: Init

```

1  $\epsilon \leftarrow \frac{1}{n+1}$ ;
2  $\forall v \in V : h(v) \leftarrow 0$ ;
3  $f(e) \leftarrow \begin{cases} 0 & \text{if } c(e) \geq 0 \\ U(e) & \text{if } c(e) < 0 \end{cases}$ ;

```

In this way we can start with a residual graph with non-negative costs.

Definition 3.3.1 (Active and Admissible). A vertex v is **active** if $\text{ex}_f(v) > 0$.

An edge (u, v) is **admissible** if $c_{f,h}(u, v) < 0$.

Algorithm 3.2: Min Cost Flow

```

1 Run Algorithm 3.1 and report if no flow is found;
2 while  $\exists$  active  $u$  do
3   if  $\exists$  admissible edge  $(u, v)$  then
4     Push  $\min\{U_f(u, v), \text{ex}_f(u)\}$  units of flow from  $u$  to  $v$ ;
5   else
6      $h(u) \leftarrow h(u) + \epsilon$ ;
7 return  $f$ ;

```

3.3.1 Correctness

First of all we claim some of the invariants of the algorithm.

Lemma 3.3.1.

- f is integral and respecting capacity. There're no new deficit vertices introduced.
- (f, h) is ε -optimal.
 - Initially, $c_{f,h}(a) \geq 0$ for all $a \in E_f$.
 - When we relabel a vertex u , there's no admissible edges incident to it, so $\forall (u, v) \in E_f : c_{f,h}(u, v) \geq 0$. After a relabel, the new reduced cost will be at least $-\varepsilon$.
 - When we push an edge, we always push along an edge that has a negative reduced cost, and bring a reverse edge with positive reduced cost. So the new reduced costs will still be at least $-\varepsilon$.
- Any deficit vertex v has $h(v) = 0$.

Then it suffices to show that the height is upper-bounded (and thus the algorithm will halt).

Lemma 3.3.2. The maximum height of any vertex is at most $n \cdot (C_{\max} + 1)$, where $C_{\max} = \max_{e \in E} |c(e)|$.

Proof. First we claim there must be a path from any excess vertex u to a deficit vertex v . Otherwise, assume a set S that contains all reachable vertices from u in G_f . Then

$$\begin{aligned} \text{ex}_f(S) &= b(S) - f^{\text{out}}(S) + f^{\text{in}}(S) \\ &= b(S) - U(\delta^{\text{out}}(S)) + 0 \\ &\leq 0 \end{aligned}$$

which contradicts the fact that $\text{ex}_f(u) > 0 \Rightarrow \text{ex}_f(S) > 0$.

Now we fix such a path P from u to v . Then we have

$$\begin{aligned} -\varepsilon \cdot |P| &\leq c_{f,h}(P) \\ &= c_f(P) - h(u) + h(v) \\ \Rightarrow h(u) &\leq c_f(P) + \varepsilon \cdot |P| \\ &\leq |P| \cdot C_{\max} + |P| \cdot \varepsilon \\ &\leq n(C_{\max} + \varepsilon) \leq n(C_{\max} + 1) \end{aligned}$$

■

Corollary 3.3.1. Every vertex is relabeled at most $L \leq n(C_{\max} + 1)/\varepsilon = O(n^2 C_{\max})$ times.

3.3.2 Time

Similar to the previous chapter, there're at most nL relabels, mL saturated pushes and mnL unsaturated pushes. Each of them takes $O(1)$ time, so the total time is $O(mnL) = O(mn^3 C_{\max})$.

Appendix

Appendix A

Algebra

Theorem A.0.1. For any polynomial $A \in \mathbb{C}[X]$ of degree $n - 1$, any point set $\{p_0, \dots, p_{n-1}\}$ of size n , $\{(p_i, A(p_i))\}_{i=0}^{n-1}$ determines $A(x)$ and vice versa.

Theorem A.0.2 (Fundamental Theorem of Algebra). Every non-constant single-variable polynomial with complex coefficients has at least one complex root.

In math terms, for $A(x) = \sum_{i=0}^{n-1} a_i x^i$, either $\exists x \in \mathbb{C}$ such that $A(x) = 0$, or $a_0 = a_1 = \dots = a_{n-1} = 0$.

Appendix B

Analysis

Theorem B.0.1 (Master Theorem (simplified)). Let the recurrence relation be:

$$T(n) = aT\left(\frac{n}{b}\right) + n^d$$

where $a \geq 1$ is the number of subproblems, $b > 1$ is the factor by which the problem size is reduced, and $d \geq 0$. The asymptotic bound of $T(n)$ is given by:

1. If $\log_b a > d$, then $T(n) = O(n^{\log_b a})$.
2. If $\log_b a = d$, then $T(n) = O(n^d \log n)$.
3. If $\log_b a < d$, then $T(n) = O(n^d)$.

Theorem B.0.2 (Master Theorem). Let $a \geq 1$ and $b > 1$ be constants, let $f(n)$ be a function, and let $T(n)$ be defined on the nonnegative integers by the recurrence:

$$T(n) = aT(n/b) + f(n)$$

Then $T(n)$ has the following asymptotic bounds:

- **Case 1:** If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$.
- **Case 2:** If $f(n) = \Theta(n^{\log_b a} \log^k n)$ for some constant $k \geq 0$, then $T(n) = \Theta(n^{\log_b a} \log^{k+1} n)$.
- **Case 3:** If $f(n) = \Omega(n^{\log_b a + \epsilon})$ for some constant $\epsilon > 0$, and if $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n (Regularity Condition), then $T(n) = \Theta(f(n))$.